

ORIE5550_HW4_Markdown

Luis Alonso Cendra Villalobos (lc2234)


2024-02-28

Question 2

Consider the time series data `gnp` from the R package `astsa`.

(a) Take a suitable preliminary transformation of the series, and produce its time plot; In the following parts, work with the transformed series.

```
gnp_data <- ts(gnp, start = c(1947, 1), end = c(2002, 3) , frequency = 4)
gnp_data = log(gnp_data)
plot(gnp_data, type="l", main="Box-Cox Log-Transform of the GNP data",
     ylab = "log(gnp)", xlab = "Year")
```



(b) Leave out the last 5 observations. Denote these samples as test data. The remaining observations are your training data. Fit a quadratic trend to the series using regression with the training data; Produce a time plot and a correlogram of the residuals obtained after removing the trend from the series

```
# Leave out the last 5 observations as test data
training_data <- head(gnp_data, -5)
test_data <- tail(gnp_data, 5)

tt <- seq(1, length(training_data), by=1)
tt2 <- tt^2


fitModel <- lm(training_data ~ tt + tt2)
summary(fitModel)

##
## Call:
## lm(formula = training_data ~ tt + tt2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.073915 -0.021840 -0.000729  0.024088  0.057608
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  7.289e+00  6.264e-03 1163.78  <2e-16 ***
## tt          1.025e-02  1.321e-04   77.63  <2e-16 ***
## tt2         -8.755e-06  5.840e-07  -14.99  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.03054 on 215 degrees of freedom
## Multiple R-squared:  0.9967, Adjusted R-squared:  0.9966
## F-statistic: 3.225e+04 on 2 and 215 DF,  p-value: < 2.2e-16


residuals <- fitModel$residuals

par(mfrow = c(1, 2))
plot(residuals, type="l", main="Residuals of the fitted model")
acf(residuals, main="ACF of Residuals of the fitted model")
```

Residuals of the fitted model



ACF of Residuals of the fitted model



(c) Fit an ARMA(p,q) model to the residual of the regression with an order p,q determined by an information criterion; Include the output; Produce the sample ACF and PACF of the residuals of the ARMA(p,q) model; Check the assumptions of white noise and normality for the residuals

```
auto.arima(residuals,max.p=10,max.q=10,ic="aic",allowmean = FALSE) # AIC
```

```
## Series: residuals
## ARIMA(1,0,3) with zero mean
##
## Coefficients:
##          ar1      ma1      ma2      ma3
##          0.8499  0.3880  0.3237  0.1553
## s.e.      0.0430  0.0744  0.0774  0.0733
##
## sigma^2 = 8.49e-05: log likelihood = 713.16
## AIC=-1416.32   AICc=-1416.03   BIC=-1399.39
```


```
auto.arima(residuals,max.p=10,max.q=10,ic="bic",allowmean = FALSE) # BIC
```

```
## Series: residuals
## ARIMA(2,0,0) with zero mean
##
```

```
## Coefficients:
##          ar1      ar2
##       1.3015 -0.3812
## s.e.  0.0623  0.0624
##
## sigma^2 = 8.638e-05: log likelihood = 710.3
## AIC=-1414.61   AICc=-1414.5   BIC=-1404.46

arma.model <- arima(residuals, order=c(1,0,3), include.mean = FALSE, method = "ML")
par(mfrow = c(1, 2))
acf(arma.model$residuals, main="ACF of Residuals from the ARMA model")
pacf(arma.model$residuals, lag.max=20,
     ylim=c(-1, 1), main = "PACF of Residuals from the ARMA model") # PACF
```

ACF of Residuals from the ARMA mPACF of Residuals from the ARMA n



```
# QQ plot
qqnorm(arma.model$residuals, main = "QQ-plot of the regression residuals")
qqline(arma.model$residuals, col = "red", lwd = 2)


# Shapiro-Wilks test. Ho is normality
shapiro.test(arma.model$residuals)

##
## Shapiro-Wilk normality test
##
## data: arma.model$residuals
## W = 0.97059, p-value = 0.0001647
```

```
#Box-Ljung test
Box.test(arma.model$residuals, lag=20, type="Ljung-Box")
```

```
##
## Box-Ljung test
##
## data: arma.model$residuals
## X-squared = 13.588, df = 20, p-value = 0.8508
```

QQ-plot of the regression residuals



We favor the ARIMA(1,0,3) due to its higher log-likelihood and lower information criteria.

From the Box-Ljung test, we cannot reject the no-autocorrelation among residuals up to lag 20, in accordance to the assumptions of white noise. The QQ-plot of the residuals exhibit heavy tails and we cannot reject the null hypothesis of normality for the Shapiro-Wilks test, thus, evidence suggests our residuals are not normally distributed; nevertheless, this is not a violation to white noise assumptions as there's no particular requirement that these are normally distributed.

(d) Forecast the transformed time series for 5 steps into the future; Compute the mean squared forecast error (MSFE) by using the test data; Provide the outputs.

```
arma.model <- arima(residuals, order=c(1,0,3), include.mean = FALSE, method = "ML")

h = 5
arma.forecast <- predict(arma.model, h)
round(arma.forecast$pred, 3)
```

```
## Time Series:
## Start = 219
## End = 223
## Frequency = 1
## [1] 0.016 0.012 0.010 0.008 0.007
```

```
round(arma.forecast$se, 3)
```

```
## Time Series:
## Start = 219
## End = 223
## Frequency = 1
## [1] 0.009 0.015 0.019 0.023 0.025
```

```
forecasted_residuals <- matrix(0, nrow = 1, ncol = 5)
forecasted_residuals[1,1] <- as.numeric(arma.forecast$pred)[1]
forecasted_residuals[1,2] <- as.numeric(arma.forecast$pred)[2]
forecasted_residuals[1,3] <- as.numeric(arma.forecast$pred)[3]
forecasted_residuals[1,4] <- as.numeric(arma.forecast$pred)[4]
forecasted_residuals[1,5] <- as.numeric(arma.forecast$pred)[5]
```

```
estimated_coeffs <- matrix(0, nrow= 1, ncol = 3)
estimated_coeffs[1,1] <- fitModel$coefficients[1]
estimated_coeffs[1,2] <- fitModel$coefficients[2]
estimated_coeffs[1,3] <- fitModel$coefficients[3]
```

```
time_matrix <- matrix(0, nrow = 5, ncol = 3)
time_matrix[, 1] <- 1
time_matrix[, 2] <- 219:223
time_matrix[, 3] <- time_matrix[, 2]^2
time_matrix <- t(time_matrix)
```

```
trend.forecast = estimated_coeffs %*% time_matrix
ts.forecast.pred <- trend.forecast + forecasted_residuals
ts.forecast.pred
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] 9.130672 9.133126 9.137347 9.142227 9.147312
```

```
lower <- ts.forecast.pred - (qnorm(0.975) * forecasted_residuals)
upper <- ts.forecast.pred + (qnorm(0.975) * forecasted_residuals)
```

```
msfe <- mean( (ts.forecast.pred - as.numeric(test_data))^2 )
msfe
```

```
## [1] 3.667099e-05
```

(e) Now, fit an ARIMA(p,d,q) model to the training data; Forecast the transformed time series for 5 steps into the future by using this model; Compute the MSFE again; Which model is preferred in terms of minimal MSFE? Include the output

```
auto.arima(training_data,max.p=10,max.q=10, max.d=2, ic="aic",allowmean = FALSE) # AIC
```

```
## Series: training_data
## ARIMA(1,1,0) with drift
##
## Coefficients:
##          ar1    drift
##          0.3473 0.0084
## s.e.    0.0636 0.0010
##
## sigma^2 = 9.275e-05: log likelihood = 700.81
## AIC=-1395.63   AICc=-1395.51   BIC=-1385.49
```

```
auto.arima(training_data,max.p=10,max.q=10, max.d=2, ic="bic",allowmean = FALSE) # BIC
```

```
## Series: training_data
## ARIMA(1,1,0) with drift
##
## Coefficients:
##          ar1    drift
##          0.3473 0.0084
## s.e.    0.0636 0.0010
##
## sigma^2 = 9.275e-05: log likelihood = 700.81
## AIC=-1395.63   AICc=-1395.51   BIC=-1385.49
```

```
arma.model.2e <- arima(training_data, order=c(1,1,0), include.mean = FALSE, method = "ML")
```

```
h = 5
arma.forecast.2e <- predict(arma.model.2e, h)
round(arma.forecast.2e$pred, 3)
```

```
##          Qtr1  Qtr2  Qtr3  Qtr4
## 2001                9.128 9.128
## 2002 9.127 9.127 9.127
```

```
msfe <- mean( (arma.forecast.2e$pred - as.numeric(test_data))^2 )
msfe
```

```
## [1] 0.0003259941
```

We can see that the ARMA(p,q) model minimizes the MSFE better than ARMA(p,d,q), for a forecasting horizon of 5 periods. This can be an unexpected result, but we can reconcile these results by recognizing the fact that the test data has a somewhat of a linear monotonically increasing trend: {9.12, 9.13, 9.14, 9.15, 9.16}. This behavior is well captured by a trend model, as opposed to data with fluctuations which we believe the ARIMA model should be better at capturing. Additionally, the test data sample is small with only 5 observations, thus, making the measurement of the models performance challenging.



Question 3

(a) Use `set.seed(99)` to generate the following random walk with drift: $X_t = -0.2t + 0.8 \sum_{s=1}^t Z_s$ with $t = 1, \dots, 100$ where Z_t is IID standard normal; Include a time plot of the series and sample ACF and PACF of the series.


```
set.seed(99)
n_steps <- 100
Z_t <- rnorm(n_steps)
t <- seq_len(n_steps)
X_t <- -1 * 0.2 * t + 0.8 * cumsum(Z_t)

par(mfrow = c(2,2))
plot(t, X_t, type = 'l', col = 'blue', xlab = 'Time', ylab = 'X_t',
     main = 'Random Walk: X_t = -0.2*t + 0.8* sum(Z_t)')
acf(X_t, main="ACF of Residuals from the ARMA model")
pacf(X_t, lag.max=20, ylim=c(-1, 1), main = "PACF of Residuals from the ARMA model") # PACF
```

Random Walk: $X_t = -0.2t + 0.8 \sum(Z$ ACF of Residuals from the ARMA mod



PACF of Residuals from the ARMA mod



(b) Use the function `auto.arima` with suitable inputs to recover this model. That is, can you find an outcome of the model indicating the random walk with drift?


```
auto.arima(X_t,max.p=20,max.q=20, max.d=2, ic="aic", allowdrift=TRUE) # AIC
```

```
## Series: X_t
## ARIMA(0,1,1) with drift
##
## Coefficients:
##          ma1      drift
##          0.2996 -0.2830
## s.e.  0.1001   0.0905
##
## sigma^2 = 0.492: log likelihood = -104.4
## AIC=214.8   AICc=215.06   BIC=222.59
```

```
auto.arima(X_t,max.p=20,max.q=20, max.d=2, ic="bic", allowdrift=TRUE) # BIC
```

```
## Series: X_t
## ARIMA(0,1,1) with drift
##
## Coefficients:
##          ma1      drift
##          0.2996 -0.2830
## s.e.  0.1001   0.0905
##
## sigma^2 = 0.492: log likelihood = -104.4
## AIC=214.8   AICc=215.06   BIC=222.59
```

Thus, suggesting that the series is integrated of order 1 or random walk with a drift parameter. Given that we know the true model of the data, we can compare the estimations with the true parameters; for instance, the drift parameter is ≈ -0.28 close to the true value of -0.2 and the variance of the model is close to the true variance parameter of $(0.8)^2 \approx 0.64$.

(c) Go through the testing procedure for unit roots with significance level $\alpha = 0.05$ (for all steps). Indicate the conclusion at each step of the procedure; Check if the testing result corresponds to the used model

```
ur.gt <- ur.df(X_t, lags=0, type='trend')
summary(ur.gt)
```

```
##
## #####
## # Augmented Dickey-Fuller Test Unit Root Test #
## #####
##
## Test regression trend
##
##
## Call:
## lm(formula = z.diff ~ z.lag.1 + 1 + tt)
##
```

```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.16040 -0.42329  0.03346  0.41717  1.42773
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.545465   0.146325  -3.728 0.000327 ***
## z.lag.1      0.001874   0.025610   0.073 0.941817
## tt          0.005879   0.009311   0.631 0.529242
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7154 on 96 degrees of freedom
## Multiple R-squared:  0.04302,    Adjusted R-squared:  0.02308
## F-statistic: 2.158 on 2 and 96 DF,  p-value: 0.1212
##
##
## Value of test-statistic is: 0.0732 6.7048 2.1578
##
## Critical values for test statistics:
##      1pct  5pct 10pct
## tau3 -4.04 -3.45 -3.15
## phi2  6.50  4.88  4.16
## phi3  8.73  6.49  5.47
```

We cannot reject the null $\tau : \pi = 0$ at a significance level of 5%. Therefore, we proceed to the existence of a trend by testing if $\beta_2 = 0$ given $\pi = 0$ with the null $\phi_3 : (\pi, \widetilde{\beta}_1, \widetilde{\beta}_2) = (0, \widetilde{\beta}_1, 0)$. Similarly, we cannot reject this null at a significance level of 5%. Therefore, we take $\beta_2 = 0$ and proceed to fit another model without the trend term.

```
ur.gt <- ur.df(X_t, lags=0, type='drift')
summary(ur.gt)
```

```
##
## #####
## # Augmented Dickey-Fuller Test Unit Root Test #
## #####
##
## Test regression drift
##
##
## Call:
## lm(formula = z.diff ~ z.lag.1 + 1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.1635 -0.4045  0.0486  0.4480  1.4270
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.536400   0.145167  -3.695 0.000364 ***
## z.lag.1     -0.013696   0.006899  -1.985 0.049940 *
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7132 on 97 degrees of freedom
## Multiple R-squared:  0.03904,    Adjusted R-squared:  0.02914
## F-statistic: 3.941 on 1 and 97 DF,  p-value: 0.04994
##
##
## Value of test-statistic is: -1.9853 9.9193
##
## Critical values for test statistics:
##      1pct  5pct 10pct
## tau2 -3.51 -2.89 -2.58
## phi1  6.70  4.71  3.86
```


We cannot reject the null $\tau : \pi = 0$ at a significance level of 5%. Therefore, we proceed to the existence of a drift by testing if $\beta_1 = 0$ given $\pi = 0$ with the null $\phi_1 : (\pi, \beta_1) = (0, 0)$. In this case, we reject this null at a significance level of 5% and conclude that a drift exists. We proceed to test whether $\pi = 0$, excluding the β 's. By checking the t-statistic reported, we can state that π is significantly different then zero at a 5% confidence level, and conclude there is not a unit root.

Question 4

Raotbl3, description: This dataset contains the time series used by Darryl Holden and Roger Perman in their article: “Unit Roots and Cointegration for the Economist”

```
data(Raotbl3)
attach(Raotbl3)
lc_TS <- ts(lc, start = c(1966, 4), end = c(1991, 2), frequency = 4)
plot(lc_TS, type="l", main="Real consumption exp., U.K., 1966-Q4 to 1991-Q2.",
     ylab = "Expenditure", xlab = "Quarter")
```

Real consumption exp., U.K., 1966–Q4 to 1991–Q2.



(a) Consider the time series `Raotbl3$lc` of real consumption expenditure from the United Kingdom starting in 1966:4 until 1991:2 in the R package `urca`. Produce a time plot of the series; Go through the testing procedure for unit roots discussed in class taking $k = 3$ for the number of lagged series differences to include in the regression; Indicate the conclusion at each step of the procedure.

```
ur.gt <- ur.df(lc_TS, lags=3, type='trend')
summary(ur.gt)
```

```
##
## #####
## # Augmented Dickey-Fuller Test Unit Root Test #
## #####
##
## Test regression trend
##
##
## Call:
## lm(formula = z.diff ~ z.lag.1 + 1 + tt + z.diff.lag)
##
## Residuals:
```

| | Min | 1Q | Median | 3Q | Max |
|----|-----------|-----------|----------|----------|----------|
| ## | -0.044714 | -0.006525 | 0.000129 | 0.006225 | 0.045353 |

```
##
```

```
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.7976591  0.3547775   2.248  0.0270 *
## z.lag.1      -0.0758706  0.0338880  -2.239  0.0277 *
## tt           0.0004915  0.0002159   2.277  0.0252 *
## z.diff.lag1 -0.1063957  0.1006744  -1.057  0.2934
## z.diff.lag2  0.2011373  0.1012373   1.987  0.0500 .
## z.diff.lag3  0.2998586  0.1020548   2.938  0.0042 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01307 on 89 degrees of freedom
## Multiple R-squared:  0.1472, Adjusted R-squared:  0.09924
## F-statistic: 3.071 on 5 and 89 DF,  p-value: 0.01325
##
##
## Value of test-statistic is: -2.2389 3.7382 2.5972
##
## Critical values for test statistics:
##      1pct  5pct 10pct
## tau3 -4.04 -3.45 -3.15
## phi2  6.50  4.88  4.16
## phi3  8.73  6.49  5.47
```

We cannot reject the null $\tau : \pi = 0$ at a significance level of 5%. Therefore, we proceed to the existence of a trend by testing if $\beta_2 = 0$ given $\pi = 0$ with the null $\phi_3 : (\pi, \widetilde{\beta}_1, \widetilde{\beta}_2) = (0, \widetilde{\beta}_1, 0)$. Similarly, we cannot reject this null at a significance level of 5%. Therefore, we take $\beta_2 = 0$ and proceed to fit another model without the trend term.

```
ur.gt <- ur.df(lc_TS, lags=3, type='drift')
summary(ur.gt)
```

```
##
## #####
## # Augmented Dickey-Fuller Test Unit Root Test #
## #####
##
## Test regression drift
##
##
## Call:
## lm(formula = z.diff ~ z.lag.1 + 1 + z.diff.lag)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.047547 -0.007071  0.000265  0.007731  0.046880
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.0123237  0.0851358   0.145  0.8852
## z.lag.1      -0.0007356  0.0079043  -0.093  0.9261
## z.diff.lag1 -0.1433015  0.1016454  -1.410  0.1620
## z.diff.lag2  0.1615256  0.1020242   1.583  0.1169
```

```
## z.diff.lag3 0.2585280 0.1027364 2.516 0.0136 *
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01337 on 90 degrees of freedom
## Multiple R-squared: 0.09747, Adjusted R-squared: 0.05735
## F-statistic: 2.43 on 4 and 90 DF, p-value: 0.05335
##
##
## Value of test-statistic is: -0.0931 2.8806
##
## Critical values for test statistics:
##      1pct 5pct 10pct
## tau2 -3.51 -2.89 -2.58
## phi1 6.70 4.71 3.86
```

We cannot reject the null $\tau : \pi = 0$ at a significance level of 5%. Therefore, we proceed to the existence of a drift by testing if $\beta_1 = 0$ given $\pi = 0$ with the null $\phi_1 : (\pi, \beta_1) = (0, 0)$. Similarly, we cannot reject this null at a significance level of 5%. Therefore, we take $\beta_1 = 0$ and proceed to fit another model without trend nor drift terms.

```
ur.gt <- ur.df(lc_TS, lags=3, type='none')
summary(ur.gt)
```

```
##
## #####
## # Augmented Dickey-Fuller Test Unit Root Test #
## #####
##
## Test regression none
##
##
## Call:
## lm(formula = z.diff ~ z.lag.1 - 1 + z.diff.lag)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.047220 -0.007276  0.000229  0.007674  0.046921
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## z.lag.1      0.0004083  0.0001695   2.409  0.0180 *
## z.diff.lag1 -0.1444994  0.1007615  -1.434  0.1550
## z.diff.lag2  0.1599782  0.1009153   1.585  0.1164
## z.diff.lag3  0.2568572  0.1015353   2.530  0.0131 *
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.0133 on 91 degrees of freedom
## Multiple R-squared: 0.2546, Adjusted R-squared: 0.2218
## F-statistic: 7.77 on 4 and 91 DF, p-value: 1.967e-05
##
##
```

```
## Value of test-statistic is: 2.4089
##
## Critical values for test statistics:
##      1pct  5pct 10pct
## tau1 -2.6 -1.95 -1.61
```

We can reject the null $\tau : \pi = 0$ at a significance level of 5%. Therefore, we conclude there is no unit root.

(b) Repeat the testing procedure with the series in (a) with a smaller lag, $k = 2$; Check if a different choice of lag affects the conclusion.

```
ur.gt <- ur.df(lc_TS, lags=2, type='trend')
summary(ur.gt)

##
## #####
## # Augmented Dickey-Fuller Test Unit Root Test #
## #####
##
## Test regression trend
##
##
## Call:
## lm(formula = z.diff ~ z.lag.1 + 1 + tt + z.diff.lag)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.042654 -0.006960  0.000573  0.007325  0.048921
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.5900161  0.3634079   1.624   0.108
## z.lag.1      -0.0558308  0.0347023  -1.609   0.111
## tt           0.0003664  0.0002218   1.652   0.102
## z.diff.lag1 -0.0755700  0.1047334  -0.722   0.472
## z.diff.lag2  0.1555229  0.1044548   1.489   0.140
##
## Residual standard error: 0.01366 on 91 degrees of freedom
## Multiple R-squared:  0.05993,    Adjusted R-squared:  0.01861
## F-statistic:  1.45 on 4 and 91 DF,  p-value: 0.2239
##
##
## Value of test-statistic is: -1.6089 5.3138 1.3643
##
## Critical values for test statistics:
##      1pct  5pct 10pct
## tau3 -4.04 -3.45 -3.15
## phi2  6.50  4.88  4.16
## phi3  8.73  6.49  5.47
```

We cannot reject the null $\tau : \pi = 0$ at a significance level of 5%. Therefore, we proceed to the existence of a trend by testing if $\beta_2 = 0$ given $\pi = 0$ with the null $\phi_3 : (\pi, \widetilde{\beta}_1, \widetilde{\beta}_2) = (0, \widetilde{\beta}_1, 0)$. Similarly, we cannot reject

this null at a significance level of 5%. Therefore, we take $\beta_2 = 0$ and proceed to fit another model without the trend term.

```
ur.gt <- ur.df(lc_TS, lags=2, type='drift')
summary(ur.gt)

##
## #####
## # Augmented Dickey-Fuller Test Unit Root Test #
## #####
##
## Test regression drift
##
##
## Call:
## lm(formula = z.diff ~ z.lag.1 + 1 + z.diff.lag)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.044303 -0.007044  0.001045  0.007454  0.049709
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  6.412e-03  8.590e-02   0.075   0.941
## z.lag.1      -1.105e-05  7.969e-03  -0.001   0.999
## z.diff.lag1  -1.069e-01  1.040e-01  -1.029   0.306
## z.diff.lag2   1.299e-01  1.043e-01   1.246   0.216
##
## Residual standard error: 0.01379 on 92 degrees of freedom
## Multiple R-squared:  0.03174,    Adjusted R-squared:  0.0001708
## F-statistic: 1.005 on 3 and 92 DF,  p-value: 0.3941
##
##
## Value of test-statistic is: -0.0014 6.4846
##
## Critical values for test statistics:
##      1pct  5pct 10pct
## tau2 -3.51 -2.89 -2.58
## phi1  6.70  4.71  3.86
```

We cannot reject the null $\tau : \pi = 0$ at a significance level of 5%. Therefore, we proceed to the existence of a drift by testing if $\beta_1 = 0$ given $\pi = 0$ with the null $\phi_1 : (\pi, \widehat{\beta_1}) = (0, 0)$. We reject this null at a significance level of 5% and conclude that a drift exists. We proceed to test whether $\pi = 0$, excluding the β 's. By checking the t-statistic reported, we cannot reject π is significantly different from zero, and conclude $\pi = 0$ and there is a unit root.

Therefore, there is a change in conclusions from a different lag choice.

(c) Go through the testing procedure with the transformed gnp used in Problem 1. Here, use the entire samples. Use the lag p determined for the ARMA model. If your ARMA model contains the MA part, use selectlags, contained to ur.df function; Check if the conclusion of this problem corresponds to the preference of the model (i.e. trend stationary or non-stationary) in Problem 1. (e);

```
auto.arima(gnp_data,max.p=10,max.q=10,ic="aic",allowmean = FALSE) # AIC
```

```
## Series: gnp_data
## ARIMA(1,1,0) with drift
##
## Coefficients:
##          ar1    drift
##          0.3467  0.0083
## s.e.    0.0627  0.0010
##
## sigma^2 = 9.136e-05:  log likelihood = 718.61
## AIC=-1431.22   AICc=-1431.11   BIC=-1421.01
```

```
auto.arima(gnp_data,max.p=10,max.q=10,ic="bic",allowmean = FALSE) # BIC
```

```
## Series: gnp_data
## ARIMA(1,1,0) with drift
##
## Coefficients:
##          ar1    drift
##          0.3467  0.0083
## s.e.    0.0627  0.0010
##
## sigma^2 = 9.136e-05:  log likelihood = 718.61
## AIC=-1431.22   AICc=-1431.11   BIC=-1421.01
```

For this result, we see that $p = 1$ and according to the lectures material, we select $k = p - 1 = 0$ for the ADF test.

```
ur.gt <- ur.df(gnp_data, lags=0, type='trend')
summary(ur.gt)
```

```
##
## #####
## # Augmented Dickey-Fuller Test Unit Root Test #
## #####
##
## Test regression trend
##
##
## Call:
## lm(formula = z.diff ~ z.lag.1 + 1 + tt)
##
## Residuals:
##          Min          1Q        Median          3Q          Max
```

```
## -0.037285 -0.004751 0.000541 0.005980 0.028789
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 0.2157840 0.1141539 1.890 0.0600 .
## z.lag.1     -0.0279548 0.0155053 -1.803 0.0728 .
## tt          0.0002174 0.0001292 1.682 0.0939 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01009 on 219 degrees of freedom
## Multiple R-squared: 0.02325, Adjusted R-squared: 0.01433
## F-statistic: 2.607 on 2 and 219 DF, p-value: 0.07606
##
## Value of test-statistic is: -1.8029 52.303 2.6068
##
## Critical values for test statistics:
##      1pct 5pct 10pct
## tau3 -3.99 -3.43 -3.13
## phi2  6.22  4.75  4.07
## phi3  8.43  6.49  5.47
```

We cannot reject the null $\tau : \pi = 0$ at a significance level of 5%. Therefore, we proceed to the existence of a trend by testing if $\beta_2 = 0$ given $\pi = 0$ with the null $\phi_3 : (\pi, \widetilde{\beta}_1, \widetilde{\beta}_2) = (0, \widetilde{\beta}_1, 0)$. Similarly, we cannot reject this null at a significance level of 5%. Therefore, we take $\beta_2 = 0$ and proceed to fit another model without the trend term.

```
ur.gt <- ur.df(gnp_data, lags=0, type='drift')
summary(ur.gt)
```

```
##
## #####
## # Augmented Dickey-Fuller Test Unit Root Test #
## #####
##
## Test regression drift
##
##
## Call:
## lm(formula = z.diff ~ z.lag.1 + 1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.036450 -0.005358  0.000236  0.005875  0.030495
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 0.024555 0.010570 2.323 0.0211 *
## z.lag.1     -0.001957 0.001273 -1.537 0.1256
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```

## Residual standard error: 0.01013 on 220 degrees of freedom
## Multiple R-squared:  0.01063,    Adjusted R-squared:  0.006132
## F-statistic: 2.364 on 1 and 220 DF,  p-value: 0.1256
##
##
## Value of test-statistic is: -1.5374 76.4037
##
## Critical values for test statistics:
##      1pct  5pct 10pct
## tau2 -3.46 -2.88 -2.57
## phi1  6.52  4.63  3.81

```

We cannot reject the null $\tau : \pi = 0$ at a significance level of 5%. Therefore, we proceed to the existence of a drift by testing if $\beta_1 = 0$ given $\pi = 0$ with the null $\phi_1 : (\pi, \widehat{\beta}_1) = (0, 0)$. We reject this null at a significance level of 5% and conclude that a drift exists. We proceed to test whether $\pi = 0$, excluding the β 's. By checking the t-statistic reported, we cannot reject π is significantly different from zero, and conclude $\pi = 0$ and there is a unit root. Same conclusion as in 4.b)

Our conclusion for problem 2.e. is that gnp_data looks like an ARIMA(1,1,0) with a drift. Which means that the series is integrated or order(1), thus, not stationary.

Our conclusion from 4.c. is that the $\beta_1 \neq 0$ and $pi = 0$, thus, there exists a unit root and there exists a drift. So, we conclude that the models agree that the data is not stationary and has a drift.